Scale Partners With NVIDIA to Bring Nemotron to Enterprises Learn more

scale Products Solutions Research Resources Log In Get Started

Back to Open Positions ←

DEPARTMENT
Engineering

Frontier Agents Engineer (Forward Deployed Engineering)

San Francisco, CA

Join the team shaping the future of AI at Scale.

Apply Now >

About Scale AI

Scale AI is the data foundation for AI, helping organizations build and deploy reliable production AI applications. We partner with the world's leading enterprises and government organizations to accelerate their AI transformation through frontier AI systems that solve real business problems.

Every day, we work with organizations across finance, healthcare, manufacturing, media, telecommunications, and government to build production AI agents that automate complex workflows, reason over enterprise knowledge, and operate safely at scale.

The Opportunity

The hardest part of enterprise AI isn't building a great model—it's building AI systems that reliably operate inside complex production environments.

As a Frontier Agent Engineer (Forward Deployed Engineering), you'll work directly with strategic enterprise customers to architect, integrate, deploy, and operate production AI systems. You'll combine modern software engineering, distributed systems, cloud infrastructure, and frontier AI technologies to transform cutting-edge models into reliable enterprise software.

Unlike traditional infrastructure or platform roles, you'll work at the intersection of AI and enterprise engineering. You'll design production agent architectures, integrate with customer systems, deploy AI into mission-critical workflows, and help shape the engineering best practices for the next generation of enterprise AI.

If you enjoy solving difficult systems problems, rapidly prototyping new AI capabilities, and bringing frontier AI into production, you'll fit right in.

What You'll Build

Enterprise AI Systems

- Design scalable agent architectures that combine LLMs, retrieval, memory, tools, structured knowledge, and enterprise data into reliable production workflows.
- Build robust integrations that allow AI agents to safely interact with customer systems while meeting enterprise requirements for security, governance, and compliance.
- Rapidly prototype new AI capabilities and evolve successful prototypes into production-ready systems.

AI Platform Engineering

- Build the production infrastructure that enables frontier AI research to become reliable enterprise software.
- Develop agent runtimes, orchestration frameworks, context pipelines, execution services, and tool integrations that power production AI systems.
- Engineer systems for reliability, observability, latency, scalability, retries, fallback strategies, and graceful degradation.
- Design human-in-the-loop workflows that effectively combine AI automation with expert oversight.
- Build deployment patterns that allow AI systems to evolve safely through continuous delivery and experimentation.

Production AI Quality

- Operationalize modern AI quality systems that ensure production agents remain reliable as models, prompts, and customer data evolve.
- Deploy evaluation harnesses using offline benchmarks, online experiments, golden datasets, regression suites, and LLM-as-a-Judge to detect quality regressions before they impact customers.
- Implement tracing, observability, monitoring, guardrails, grounding, and safety mechanisms that enable production AI systems to operate with confidence.
- Partner closely with Applied AI engineers to productionize new evaluation methodologies, retrieval strategies, reasoning architectures, and emerging AI capabilities.

- Rapidly evaluate newly released models, agent frameworks, evaluation methodologies, and developer tooling, determining how they can be safely adopted into production systems.

- Partner directly with enterprise customers to understand their technical infrastructure, software architecture, and operational workflows.
- Translate ambiguous customer problems into scalable production AI architectures.
- Collaborate with customer software engineers, ML engineers, platform teams, and product organizations to deploy AI into mission-critical workflows.
- Identify reusable engineering patterns that become core capabilities across multiple enterprise deployments.

Technical Leadership

- Serve as the primary technical advisor for strategic enterprise accounts.
- Lead architecture discussions spanning distributed systems, AI infrastructure, enterprise integration, and production deployment.
- Document reusable architecture patterns, deployment strategies, integration frameworks, and operational best practices.
- Work closely with Scale's product, infrastructure, and Applied AI teams to continuously improve the platform.

What Makes This Role Different

You'll work across the full lifecycle of production AI systems:

- Architecting enterprise AI systems and agent platforms
- Building integrations across cloud infrastructure and enterprise software
- Deploying AI agents into production environments
- Operationalizing evaluation frameworks, guardrails, and observability
- Running production experiments and measuring real-world business impact

- Continuously improving deployed AI systems using customer feedback and operational telemetry

Rather than supporting a single product or platform, you'll solve diverse engineering challenges across industries and use cases, rapidly building expertise across enterprise architecture, AI systems engineering, and production deployment.

Required Qualifications

- 4+ years of software engineering experience with strong fundamentals in distributed systems, data structures, algorithms, and system design.
- Strong Python programming skills with experience building production software.
- Experience building or deploying AI-powered applications using modern LLM APIs, agent frameworks, MCP, retrieval systems, or vector databases.
- Experience with cloud platforms (AWS, Azure, or GCP) and modern production infrastructure.
- Strong problem-solving skills with the ability to navigate ambiguous technical requirements and rapidly iterate toward production solutions.
- Excellent communication skills and the ability to work directly with enterprise engineering teams.

Preferred Qualifications

Enterprise AI Engineering

- Experience deploying production AI agents or autonomous systems.
- Experience designing distributed systems, APIs, orchestration services, or large-scale backend systems.
- Experience with cloud-native infrastructure, Docker, Kubernetes, Infrastructure as Code, and CI/CD.
- Experience integrating AI systems into enterprise software environments.

AI Systems

- Familiarity with modern agent architectures, retrieval systems, tool use, memory, and context engineering.

- Experience with evaluation frameworks, LLM observability, regression testing, tracing, and AI monitoring.

Customer Engineering

- Experience working directly with enterprise customers.
- Ability to translate complex technical requirements into scalable production systems.
- Strong written and verbal communication skills.
- Experience leading architecture reviews, technical workshops, or customer design sessions.

Dual Fluency

While this role initially emphasizes enterprise engineering and production deployment, every Frontier Agent Engineer develops expertise across both Forward Deployed Engineering and Applied AI.

Over time, you'll deepen your understanding of modern agent architectures, evaluation methodologies, retrieval systems, reasoning techniques, and emerging AI technologies, while continuing to build world-class production software. Our goal is to develop engineers who can design, build, evaluate, deploy, and continuously improve frontier AI systems end to end.

Compensation packages at Scale for eligible roles include base salary, equity, and benefits. The range displayed on each job posting reflects the minimum and maximum target for new hire salaries for the position and may be inclusive of several career levels at Scale; it will be determined during the interview process based on work location and additional factors, including job-related skills, experience, qualifications, interview performance, and relevant education or training. Scale employees in eligible roles are also granted equity based compensation, subject to Board of Director approval. Your recruiter can share more about the specific salary range for your preferred location during the hiring process, and confirm whether the hired role will be eligible for equity grant. You'll also receive benefits including, but not limited to: comprehensive health, dental and vision coverage, retirement benefits, a learning and development stipend, and generous PTO. Additionally, this role may be eligible for additional benefits such as a commuter stipend.

Please reference the job posting's subtitle for where this position will be located. For pay transparency purposes, the base salary range for this full-time position in the locations of San Francisco, New York, Seattle is:

$180,000—$225,000 USD

About Us:

At Scale, our mission is to develop reliable AI systems for the world's most important decisions. Our products provide the high-quality data and full-stack technologies that power the world's leading models, and help enterprises and governments build, deploy, and oversee AI applications that deliver real impact. We work closely with industry leaders like Meta, Ernst & Young, Mayo Clinic, Time Inc., the Government of Qatar, and U.S. government agencies including the Army and Air Force. We are expanding our team to accelerate the development of AI applications.

We believe that everyone should be able to bring their whole selves to work, which is why we are proud to be an inclusive and equal opportunity workplace. We are committed to equal employment opportunity regardless of race, color, ancestry, religion, sex, national origin, sexual orientation, age, citizenship, marital status, disability status, gender identity or Veteran status.

We are committed to working with and providing reasonable accommodations to applicants with physical and mental disabilities. If you need assistance and/or a reasonable accommodation in the application or recruiting process due to a disability, please contact us at accommodations@scale.com. Please see the United States Department of Labor's Know Your Rights poster for additional information.

We comply with the United States Department of Labor's Pay Transparency provision.

PLEASE NOTE: We collect, retain and use personal data for our professional business purposes, including notifying you of job opportunities that may be of interest and sharing with our affiliates. We limit the personal data we collect to that which we believe is appropriate and necessary to manage applicants' needs, provide our services, and comply with applicable laws. Any information we collect in connection with your application will be treated in accordance with our internal policies and programs designed to protect personal data. Please see our privacy policy for additional information.

Scale Is Growing. Grow With Us.

scale

Log In Get Started

most important decisions

MANAGE YOUR COOKIE PREFERENCES

COPYRIGHT © 2026 SCALE AI, INC. ALL RIGHTS RESERVED TERMS OF USE & PRIVACY POLICY